

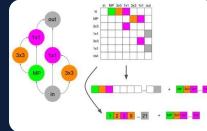
Neural Architecture Search: The Next Frontier



Colin White colin@abacus.ai

<https://abacus.ai/publications>

Research at Abacus.AI



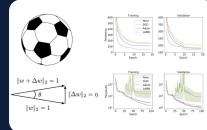
A Study on Encodings for Neural Architecture Search
(NeurIPS 2020)

NEURAL INFORMATION PROCESSING SYSTEMS
NIPS



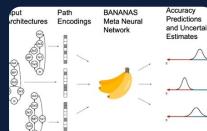
Intra-Processing Methods for Debiasing Neural Networks
(NeurIPS 2020)

NEURAL INFORMATION PROCESSING SYSTEMS
NIPS

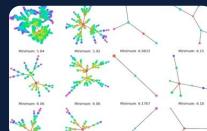


Learning by Turning: Neural Architecture Aware Optimisation
(ICML 2021)

ICML
International Conference On Machine Learning

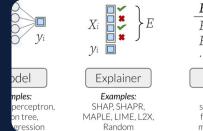


BANANAS: Bayesian Optimization with Neural Architectures for Neural Architecture Search
(AAAI 2021)



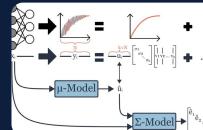
Exploring the Loss Landscape in Neural Architecture Search
(UAI 2021)

uai



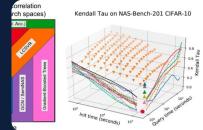
Synthetic Benchmarks for Scientific Research in Explainable Machine Learning
(ICML 2021)

ICML
International Conference On Machine Learning



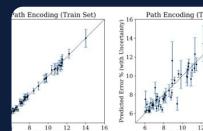
NAS-Bench-x11 and the Power of Learning Curves
(CVPR 2021 Workshop)

CVPR
VIRTUALLY JUNE 19-25



How Powerful are Performance Predictors in Neural Architecture Search?
(ICLR Workshop)

ICLR



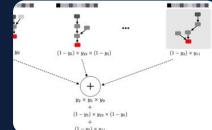
Deep Uncertainty Estimation for Model-based Neural Architecture Search
(NeurIPS 2019)

NEURAL INFORMATION PROCESSING SYSTEMS
NIPS



DECO: Debiasing through Compositional Optimization of Machine Learning Models
(NeurIPS 2019)

NEURAL INFORMATION PROCESSING SYSTEMS
NIPS



Differentiable Functions for Combining First-order Constraints with Deep Learning via Weighted Proof Tracing
(NeurIPS 2019)

NEURAL INFORMATION PROCESSING SYSTEMS
NIPS

Competitions



CVPR 2021 Unseen Data in Neural Architecture Search

2ND PLACE

<https://abacus.ai/publications>

Machine learning automation

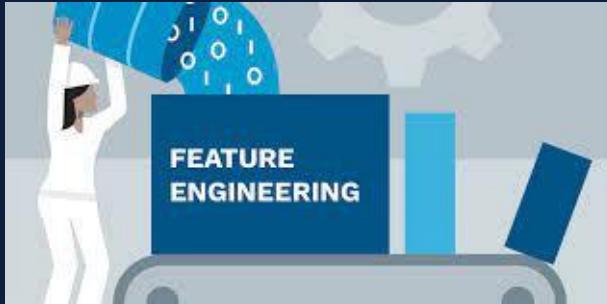


1950s

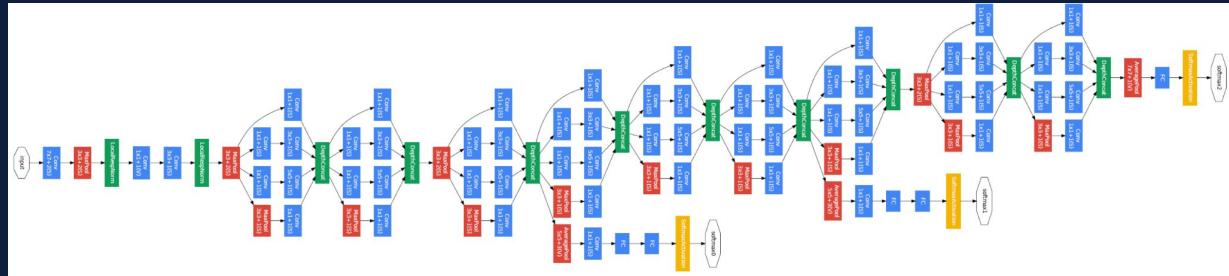
2013

2017

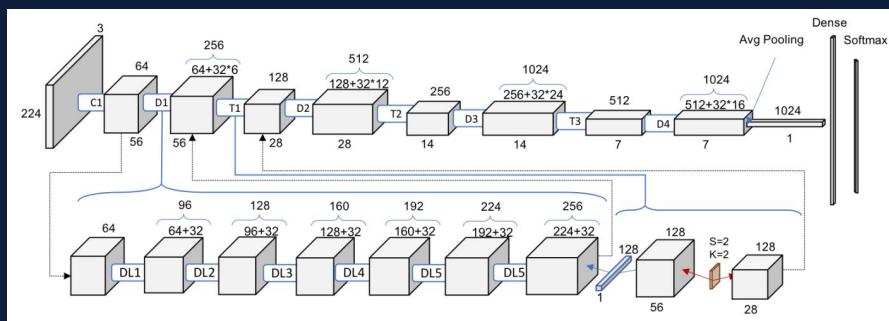
2022



Neural architecture search



[GoogLeNet \(2014\)](#)



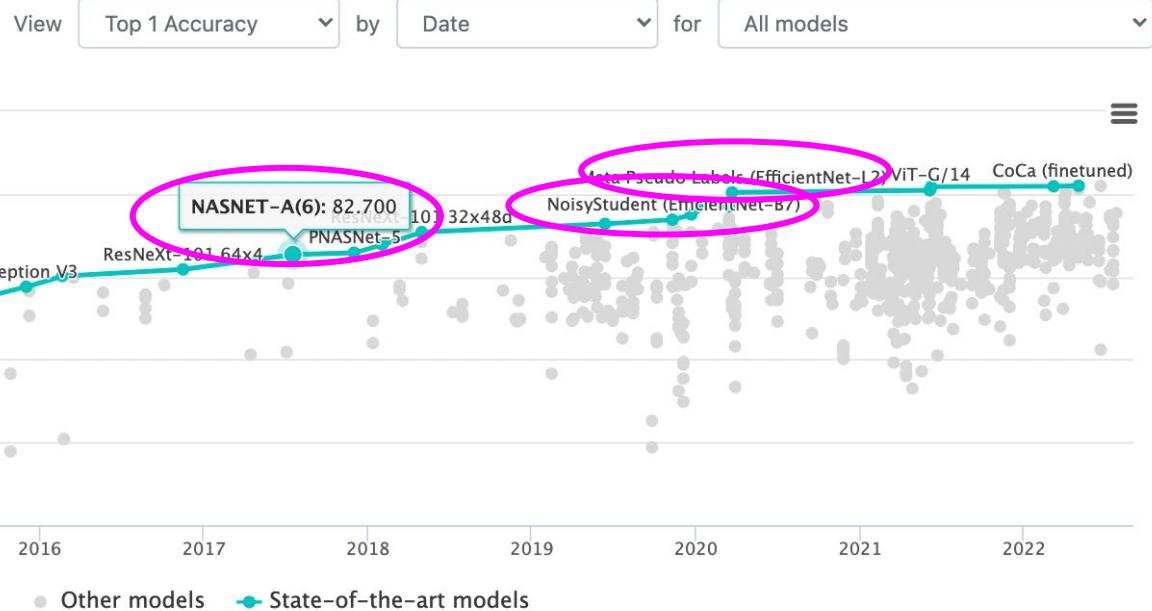
[DenseNet \(2016\)](#)

Architectures are getting increasingly more specialized and complex

Image Classification on ImageNet

Leaderboard

Dataset



What is NAS?

What is neural architecture search?

Neural architecture search (NAS) is a technique for automatically designing neural networks. NAS algorithms typically use a reinforcement learning (RL) or evolutionary algorithm (EA) to search for an optimal neural network architecture.

DALL-E My collection

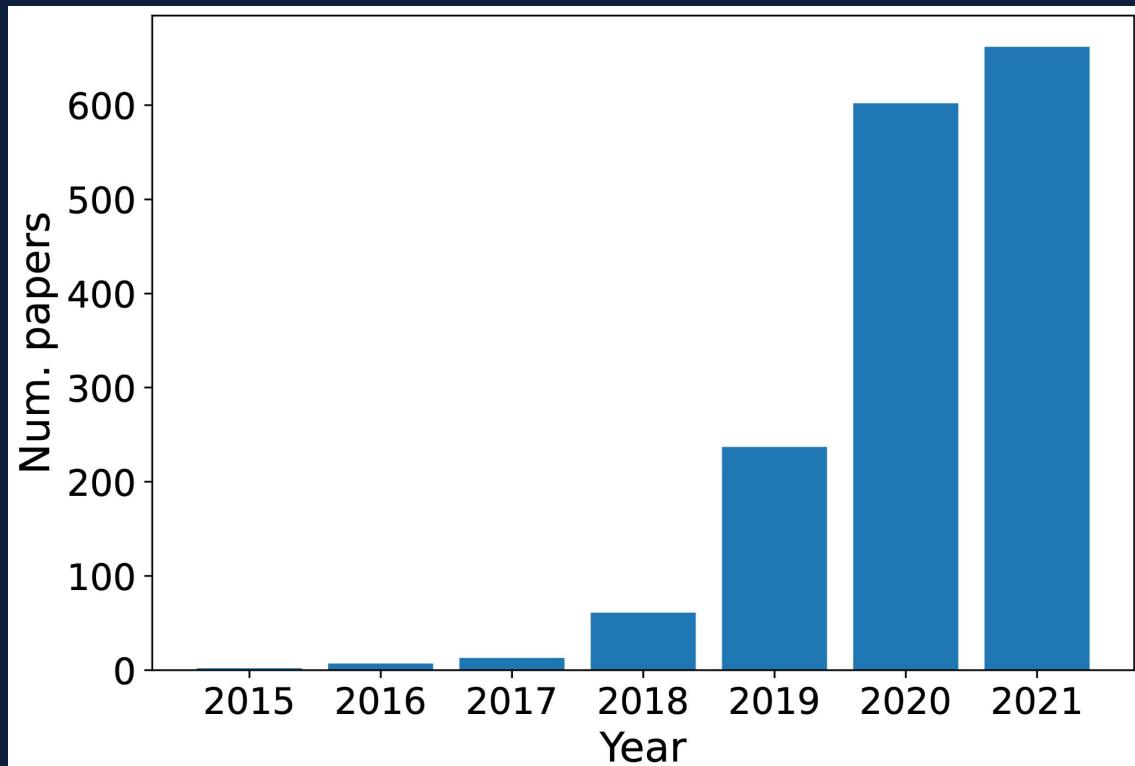
Edit the detailed description

What is neural architecture search? Generate

- A mobile application interface titled 'Autogétois' with a search bar and buttons for 'Ananaptis', 'Patsoit', and 'Sugon ñasnes'. Below it is a blue banner with the text 'Netigee Aich · Sherar Nucha?'.
- A yellow background featuring a magnifying glass over the text 'An ANarcha Adil n Aoruch' and a small orange architectural drawing.
- A white background with an orange hexagonal logo and the text 'Nutricia Atchuelthel Asudent'.
- A wooden surface background with the text 'Nur-Nchuch Acha Achutifica'.

Neural architecture search

NAS: the process of
automating the design of
neural architectures for
a given dataset.



NAS on new datasets / tasks

Spherical Omnidirectional Vision



NinaPro DB5 Prosthetics Control



FSD50K Audio Classification



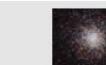
Darcy Flow PDE Solver



PSICOV Protein Folding



Cosmic Astronomy Imaging



ECG Medical Diagnostics



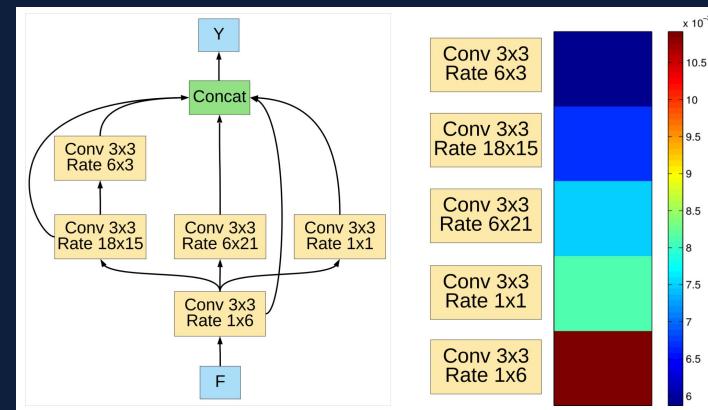
Satellite Earth Monitoring



DeepSEA Genetic Prediction



Graph neural networks
Generative adversarial network
Dense prediction tasks
Adversarial robustness
Self-supervised learning for NAS



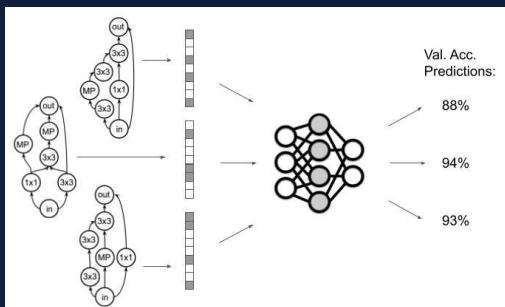
Roadmap

- Motivation and Introduction
- Performance Prediction
 - BANANAS
 - Learning curve extrapolation
 - Zero-cost proxies
- Recommender Systems
- De-biasing models

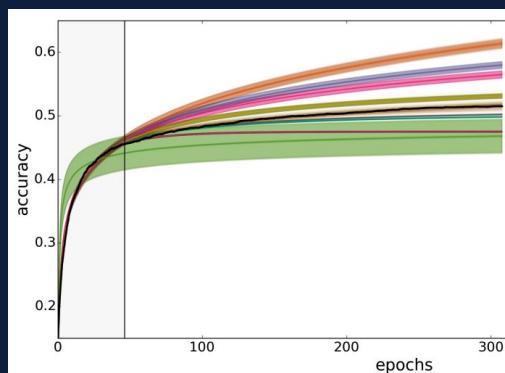


Performance Predictors

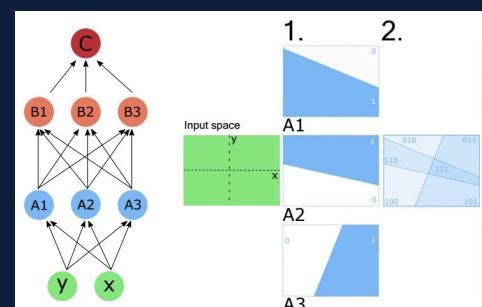
Predict the (relative) accuracy of an architecture, without fully training it.



Model-based



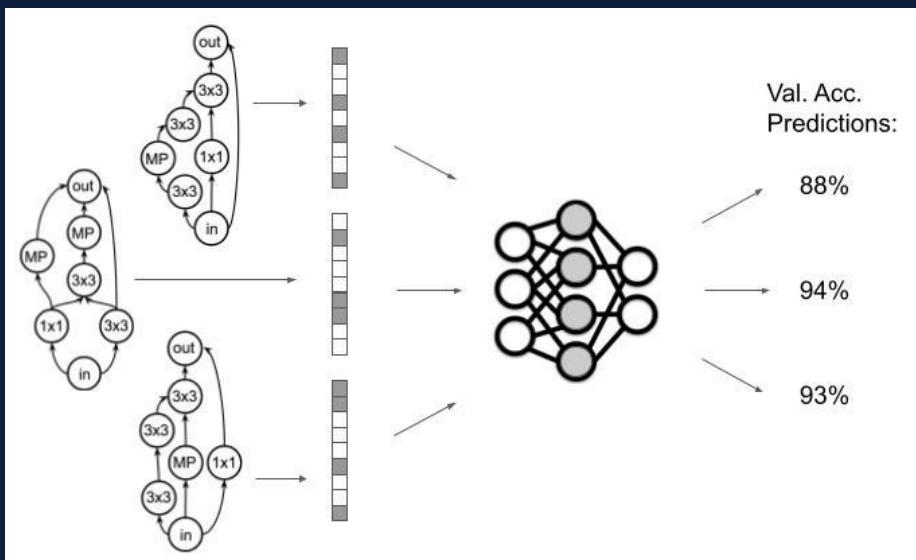
Learning curve
extrapolation



Zero-cost proxies

Model-Based Predictors

Train a surrogate model



- Gaussian processes [[Kandasamy et al. 2018](#)], [[Jin et al. 2018](#)]
- Boosted trees [[Luo et al. 2020](#)], [[Siems et al. 2020](#)]
- GNNs [[Shi et al. 2019](#)], [[Wen et al. 2019](#)]
- Specialized encodings [[White et al. 2019](#)], [[Ning et al. 2020](#)]

“BO + Neural Predictor” Framework

[NASGBO, 2019], [BONAS, 2019], [BANANAS, 2019]

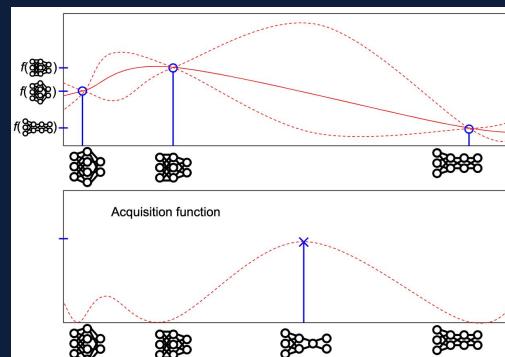
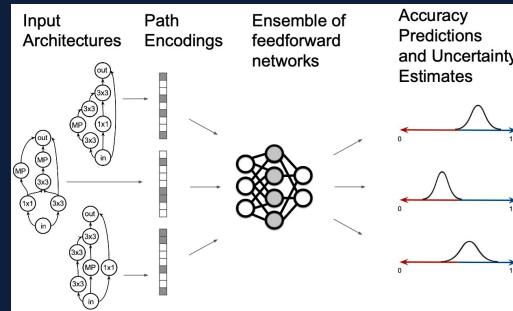
Algorithm 1 BANANAS

Input: Search space A , dataset D , parameters t_0, T, M, c, x , acquisition function ϕ , function $f(a)$ returning validation error of a after training.

1. Draw t_0 architectures a_0, \dots, a_{t_0} uniformly at random from A and train them on D .
2. For t from t_0 to T ,

- i. Train an ensemble of meta neural networks on $\{(a_0, f(a_0)), \dots, (a_t, f(a_t))\}$.
- ii. Generate a set of c candidate architectures from A by randomly mutating the x architectures a from $\{a_0, \dots, a_t\}$ that have the lowest value of $f(a)$.
- iii. For each candidate architecture a , evaluate the acquisition function $\phi(a)$.
- iv. Denote a_{t+1} as the candidate architecture with minimum $\phi(a)$, and evaluate $f(a_{t+1})$.

Output: $a^* = \operatorname{argmin}_{t=0, \dots, T} f(a_t)$.



Train 10 arch.'s each iteration

“BO + Neural Predictor” Components

Algorithm 1 BANANAS

Input: Search space A , dataset D , parameters t_0, T, M, c, x , acquisition function ϕ , function $f(a)$ returning validation error of a after training.

1. Draw t_0 architectures a_0, \dots, a_{t_0} uniformly at random from A and train them on D .
2. For t from t_0 to T ,

- i. Train an ensemble of meta neural networks on $\{(a_0, f(a_0)), \dots, (a_t, f(a_t))\}$.
 - ii. Generate a set of c candidate architectures from A by randomly mutating the x architectures a from $\{a_0, \dots, a_t\}$ that have the lowest value of $f(a)$.
 - iii. For each candidate architecture a , evaluate the acquisition function $\phi(a)$.
 - iv. Denote a_{t+1} as the candidate architecture with minimum $\phi(a)$, and evaluate $f(a_{t+1})$.
- Output:** $a^* = \operatorname{argmin}_{t=0, \dots, T} f(a_t)$.

-
- Architecture encoding
 - Uncertainty calibration
 - Neural predictor
 - Architecture
 - Acquisition optimization
 - Strategy
 - Acquisition function

BANANAS



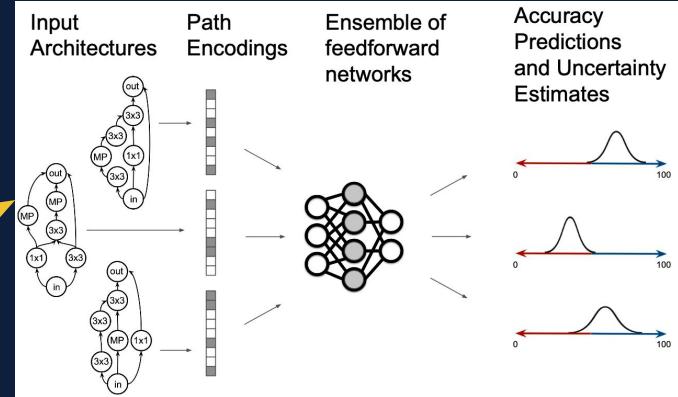
Algorithm 1 BANANAS

Input: Search space A , dataset D , parameters t_0, T, M, c, x , acquisition function ϕ , function $f(a)$ returning validation error of a after training.

1. Draw t_0 architectures a_0, \dots, a_{t_0} uniformly at random from A and train them on D .
2. For t from t_0 to T ,

- i. Train an ensemble of meta neural networks on $\{(a_0, f(a_0)), \dots, (a_t, f(a_t))\}$.
- ii. Generate a set of c candidate architectures from A by randomly mutating the x architectures a from $\{a_0, \dots, a_t\}$ that have the lowest value of $f(a)$.
- iii. For each candidate architecture a , evaluate the acquisition function $\phi(a)$.
- iv. Denote a_{t+1} as the candidate architecture with minimum $\phi(a)$, and evaluate $f(a_{t+1})$.

Output: $a^* = \operatorname{argmin}_{t=0, \dots, T} f(a_t)$.

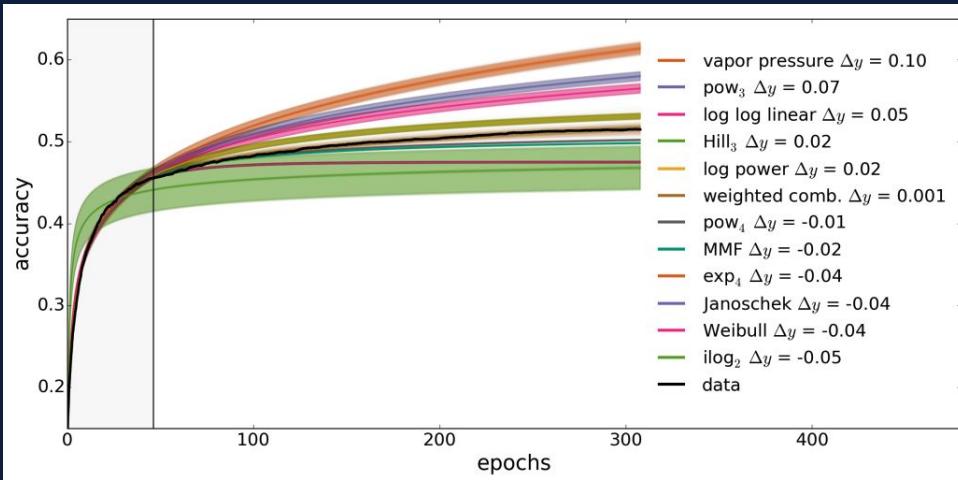


Path encoding, ensemble

Small mutations

Independent Thompson Sampling

Learning curve based predictors



- Learning curve extrapolation
 - Fit partial learning curve to parametric model [\[Domhan et al. 2015\]](#)
 - Bayesian NN [\[Klein et al. 2017\]](#)
- LCE + Surrogate
 - SVR [\[Baker et al. 2017\]](#)
 - Full LC + Bayesian NN [\[Klein et al. 2017\]](#)

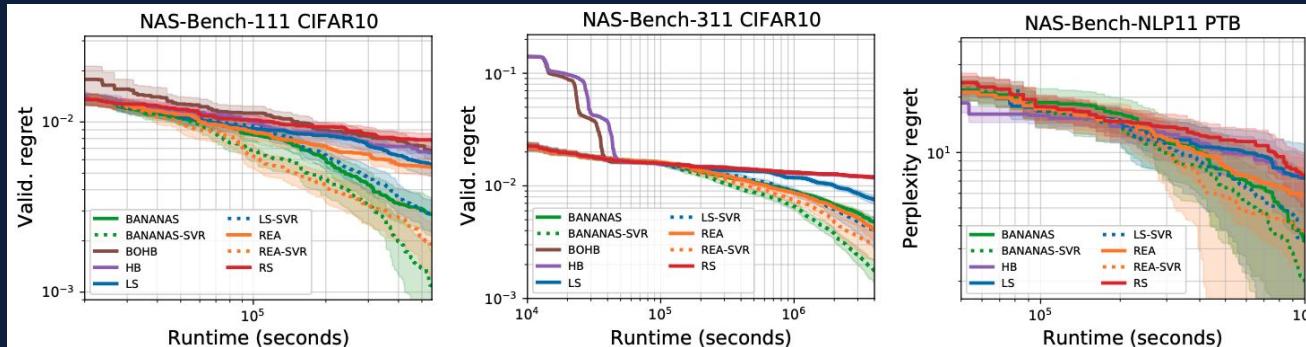
LCE Framework

Algorithm 1 Single-Fidelity Algorithm

```
1: initialize history
2: while  $t < t_{\max}$  :
3:   arches = gen_candidates(history)
4:   accs = train(arches, epoch= $E_{\max}$ )
5:   history.update(arches, accs)
6: Return arch with the highest acc
```

Algorithm 2 LCE Framework

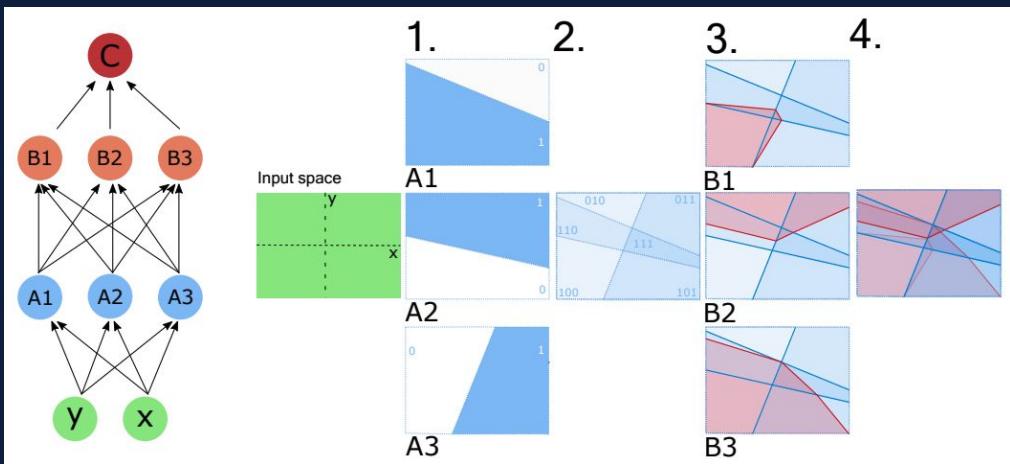
```
1: initialize history
2: while  $t < t_{\max}$  :
3:   arches = gen_candidates(history)
4:   accs = train(arches, epoch= $E_{\text{few}}$ )
5:   sorted_by_pred = LCE(arches, accs)
6:   arches = sorted_by_pred[:top_n]
7:   accs = train(arches, epoch= $E_{\max}$ )
8:   history.update(arches, accs)
9: Return arch with the highest acc
```



Zero-cost proxies

epe-nas [21]
fisher [42]
flops [25]
grad-norm [1]
grasp [43]
l2-norm [1]
jacov [23]
nwo [23]
params [25]
plain [1]
snip [14]
synflow [39]
zen-score [16]

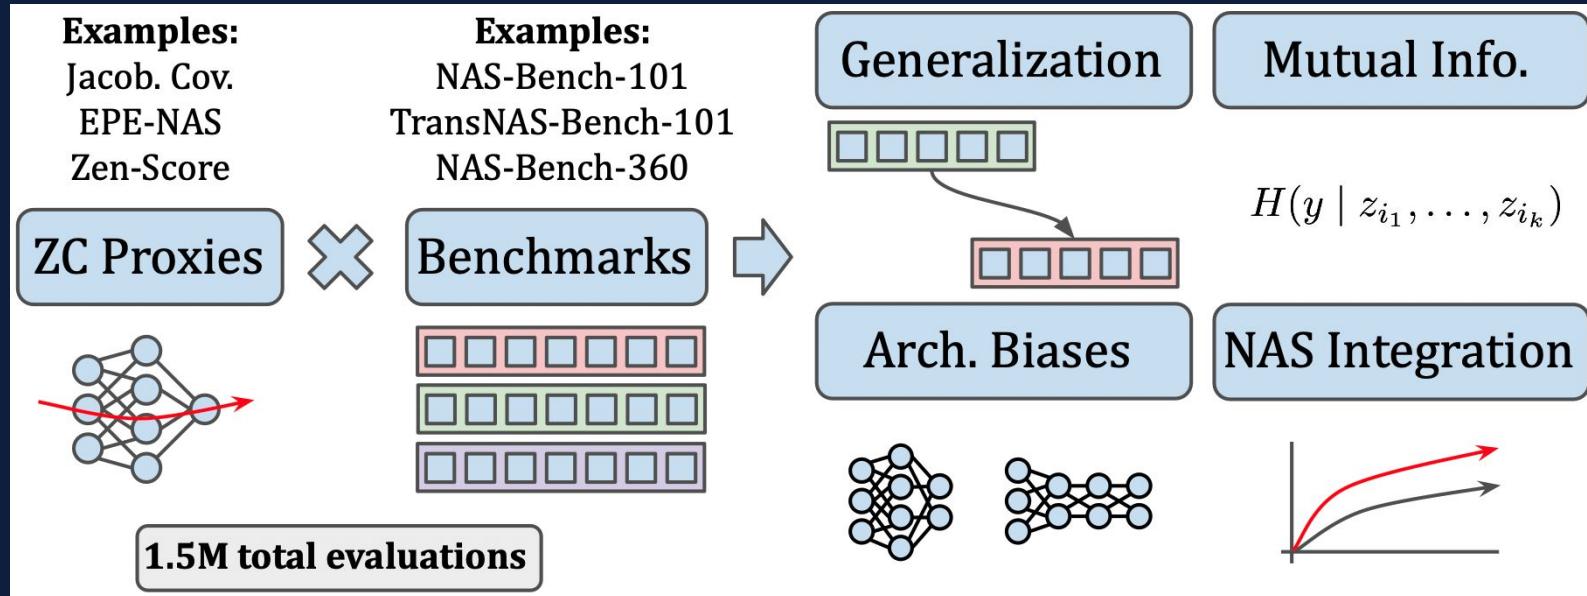
Jacobian
Pruning-at-init
Baseline
Pruning-at-init
Pruning-at-init
Baseline
Jacobian
Jacobian
Baseline
Baseline
Pruning-at-init
Pruning-at-init
Piece. Lin.



[Mellor et al. 2020]

Compute an estimate in 5 seconds

NAS-Bench-Suite-Zero (28 tasks)



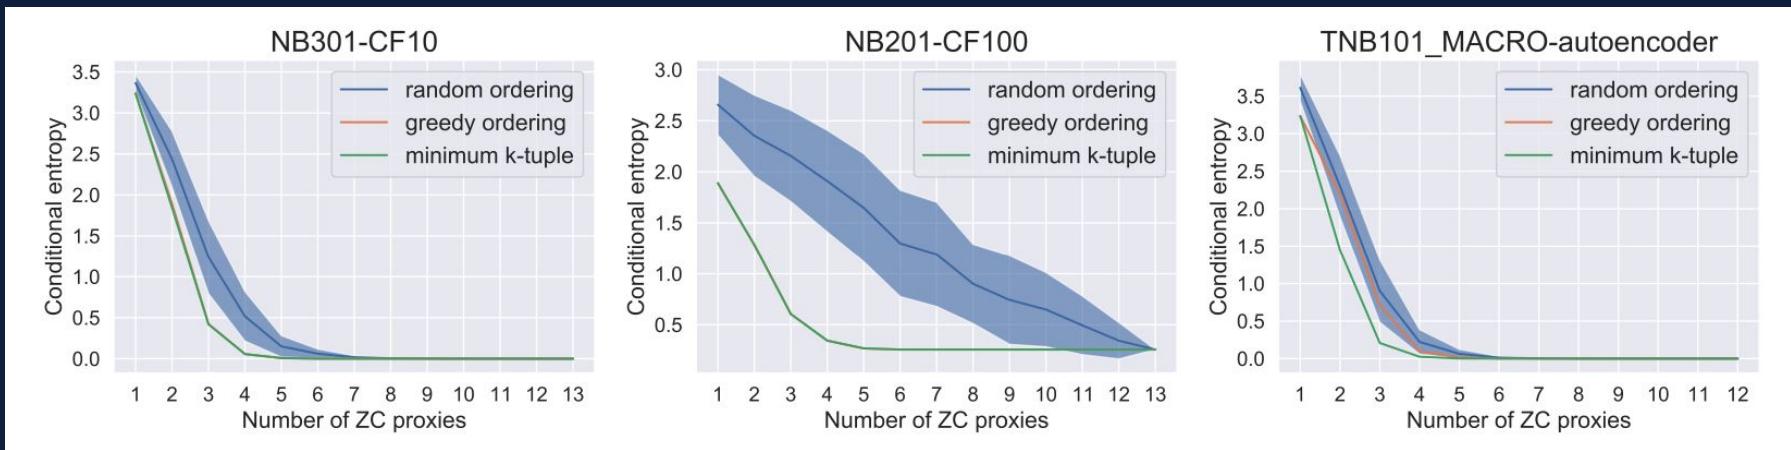
Zero-cost proxies

Table 4: Average ranking of each of the ZC proxies on each search space, and over all search spaces.

	fisher	grad_norm	grasp	jacob_cov	snip	synflow	flops	params
NATS-Bench TSS	6.0	6.0	5.0	4.0	5.67	1.33	4.0	4.0
DARTS	4.6	4.2	4.6	4.8	4.6	5.4	4.0	3.8
TransNAS-Bench-101	2.75	4.5	4.5	7.5	3.0	4.5	4.0	5.25
Overall	4.33	4.75	4.67	5.5	4.33	4.08	4.0	4.33

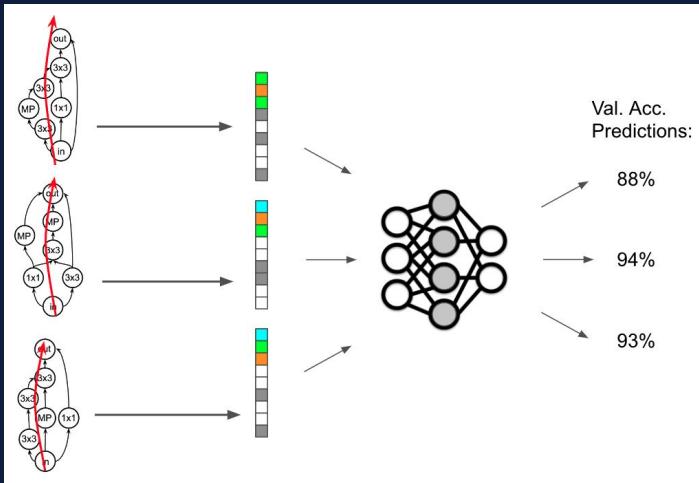
- Still do not consistently beat simple baselines
- Promising when used in conjunction with other NAS techniques

Complementary info in ZC proxies

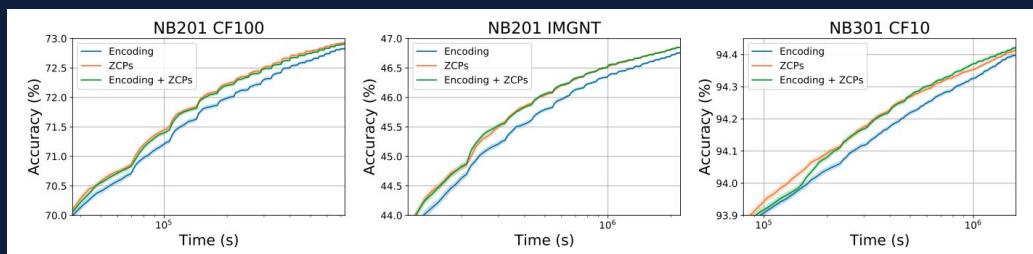


Conditional entropy $H(y \mid z_{i_1}, \dots, z_{i_k})$ vs. k

NAS integration

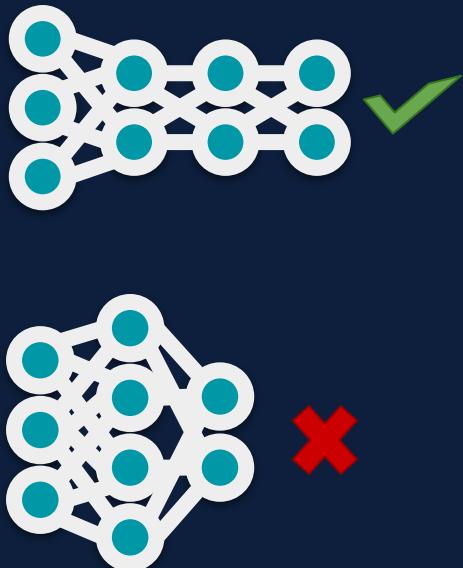


Features Benchmark	Encoding	ZC	Both	% Improvement (ZC)	% Improvement (Both)
NB101-CF10	0.546	0.708	0.718	29.67	31.50
NB201-CF10	0.622	0.905	0.906	45.50	45.66
NB201-CF100	0.640	0.907	0.908	41.71	41.87
NB201-IMGNT	0.683	0.879	0.883	28.70	29.28
NB301-CF10	0.314	0.405	0.465	28.98	48.09
TNB101_MACRO-AUTOENC	0.673	0.831	0.837	23.48	24.37
TNB101_MACRO-JIGSAW	0.809	0.706	0.809	-12.73	0.00
TNB101_MACRO-NORMAL	0.617	0.710	0.716	15.07	16.05
TNB101_MACRO-OBJECT	0.736	0.840	0.843	14.13	14.54
TNB101_MACRO-ROOM	0.683	0.589	0.707	-13.76	3.51
TNB101_MACRO-SCENE	0.832	0.891	0.899	7.09	8.05
TNB101_MACRO-SEGMENT	0.900	0.807	0.876	-10.33	-2.67
TNB101_MICRO-AUTOENC	0.714	0.754	0.803	5.60	12.46
TNB101_MICRO-JIGSAW	0.585	0.730	0.743	24.79	27.01
TNB101_MICRO-NORMAL	0.657	0.801	0.809	21.92	23.14
TNB101_MICRO-OBJECT	0.637	0.733	0.752	15.07	18.05
TNB101_MICRO-ROOM	0.582	0.843	0.844	44.85	45.02
TNB101_MICRO-SCENE	0.710	0.849	0.866	19.58	21.97
TNB101_MICRO-SEGMENT	0.767	0.886	0.897	15.51	16.95



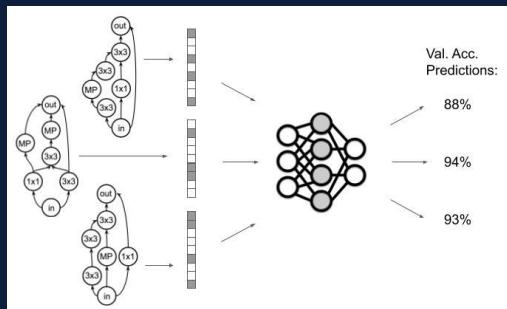
Removing biases in ZC proxies

$$f'(a) = f(a) \cdot \frac{1}{b(a) + C}$$

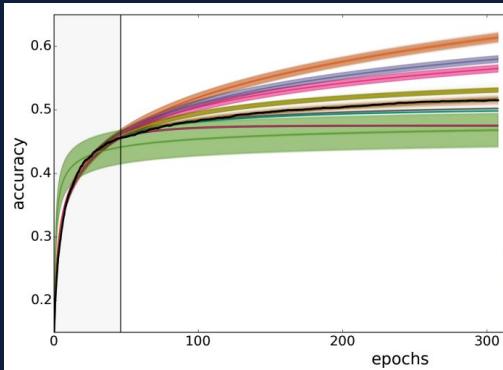


ZC proxy	dataset	bias metric	original bias	original perf.	new bias	new perf.	strategy
l2-norm	NB201-CF10	conv:pool	0.87	0.42	0.00	0.10	minimize
					0.70	0.44	equalize performance
nwot	NB301-CF10	conv:pool	0.78	0.49	0.00	0.03	minimize
					0.78	0.49	equalize performance
synflow	NB201-CF100	cell size	0.57	0.68	0.01	0.64	minimize
					0.35	0.71	equalize performance
synflow	NB201-IM	cell size	0.58	0.76	0.01	0.62	minimize
					0.46	0.76	equalize performance
flops	NB301-CF10	num. skip	-0.35	0.43	-0.01	0.06	minimize
					-0.35	0.43	equalize performance

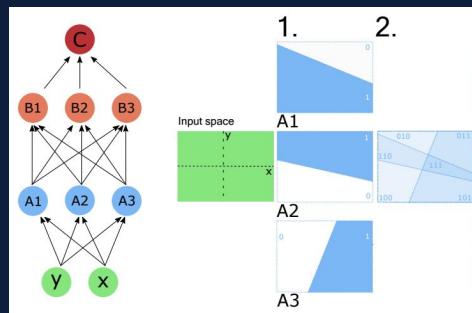
Performance predictor families



Model-based

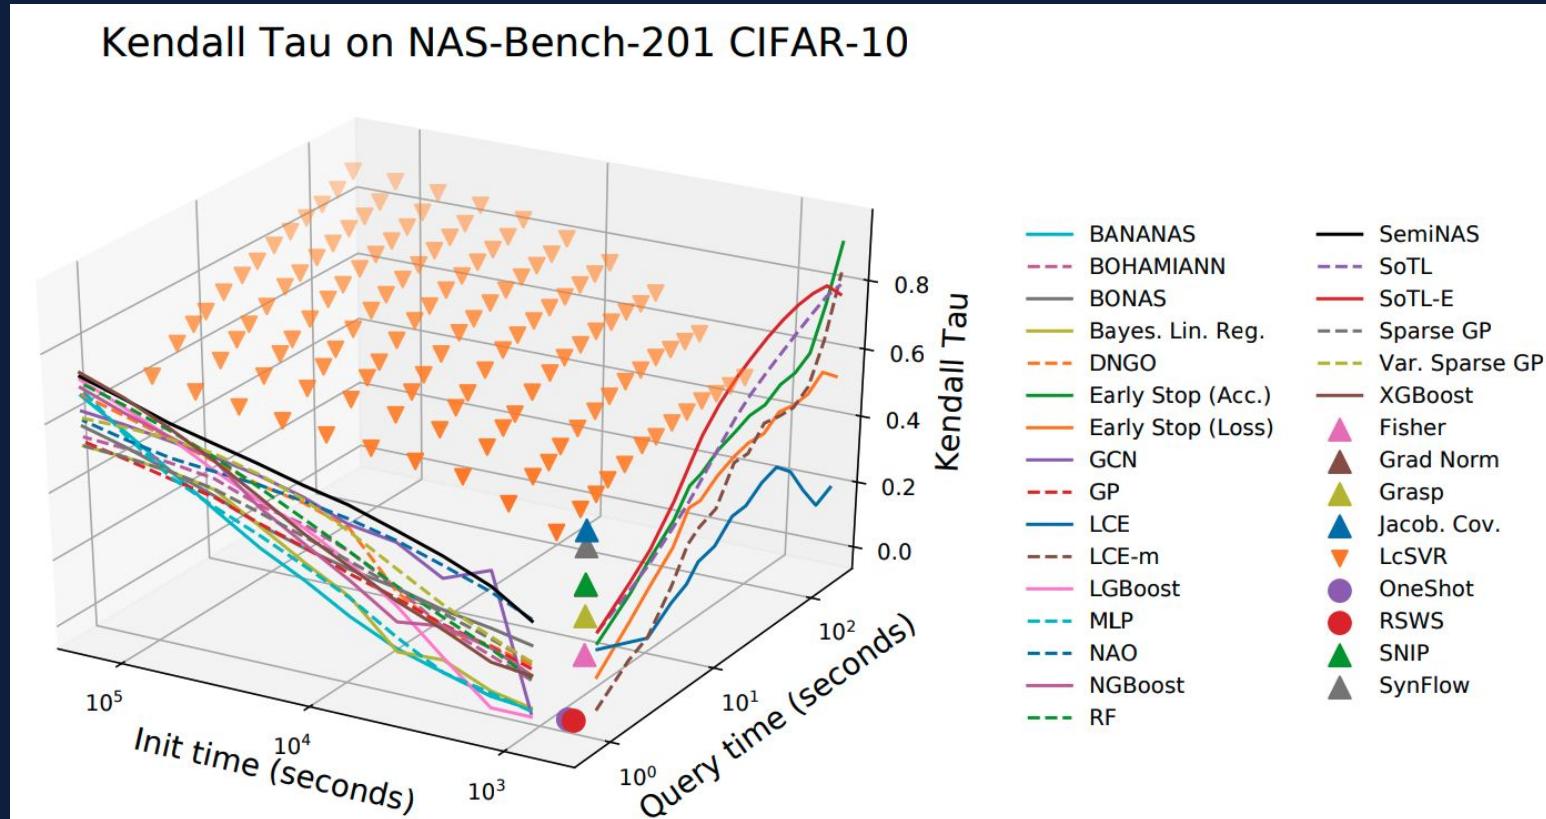


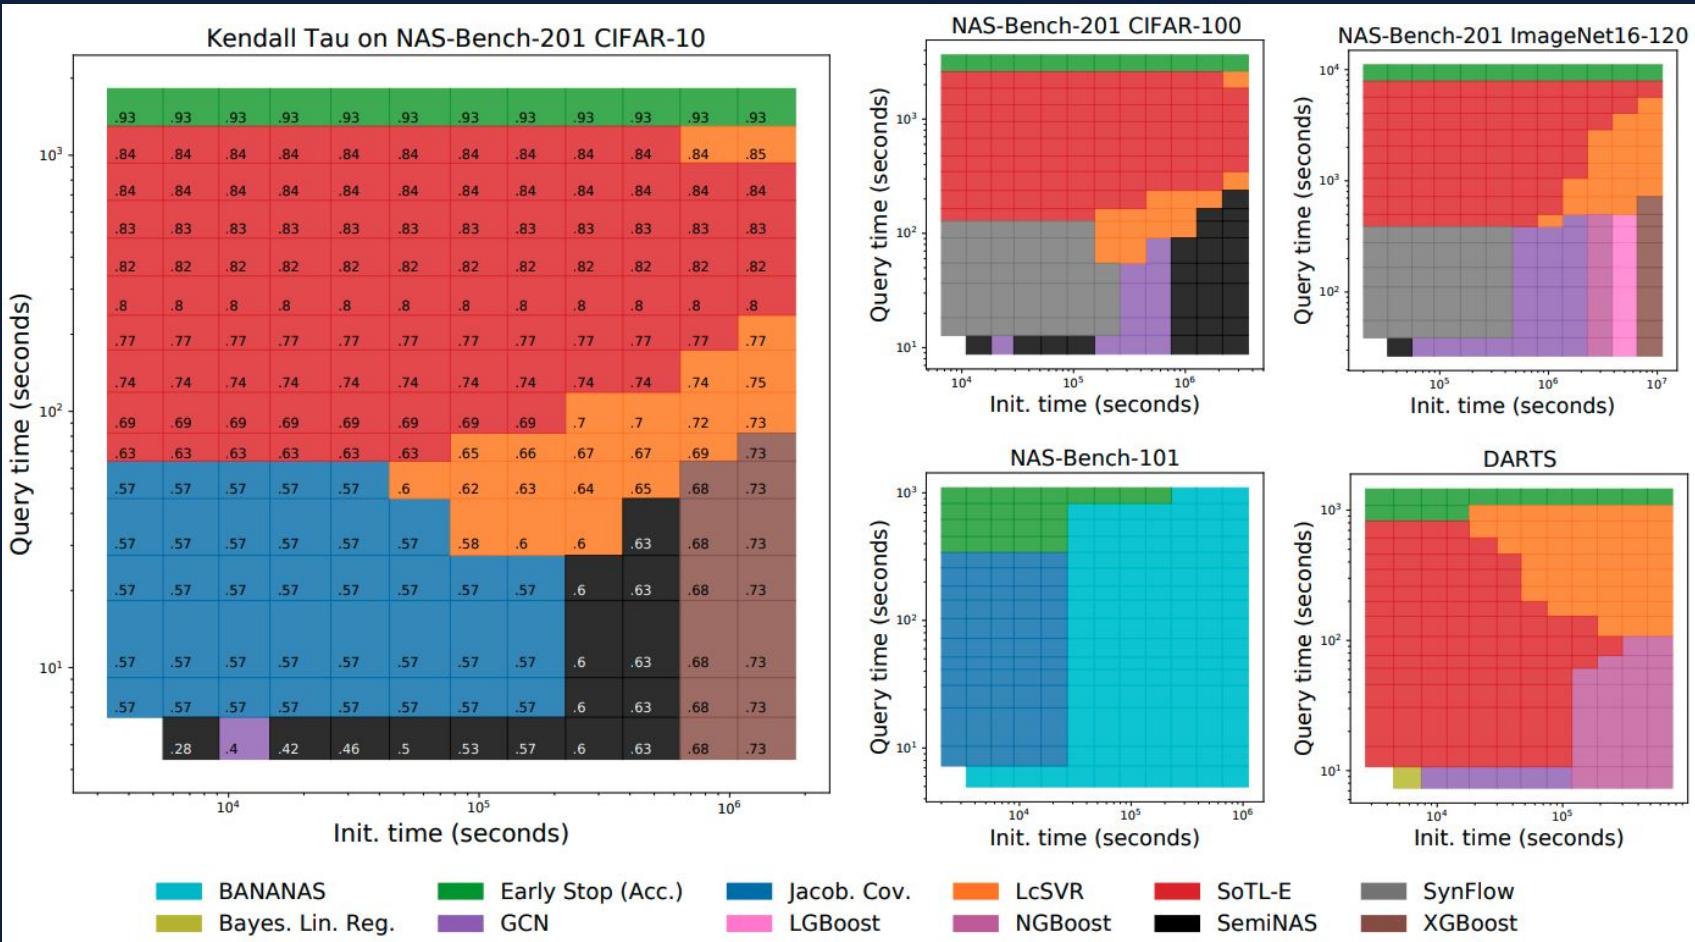
Learning curve
extrapolation



Zero-cost proxies

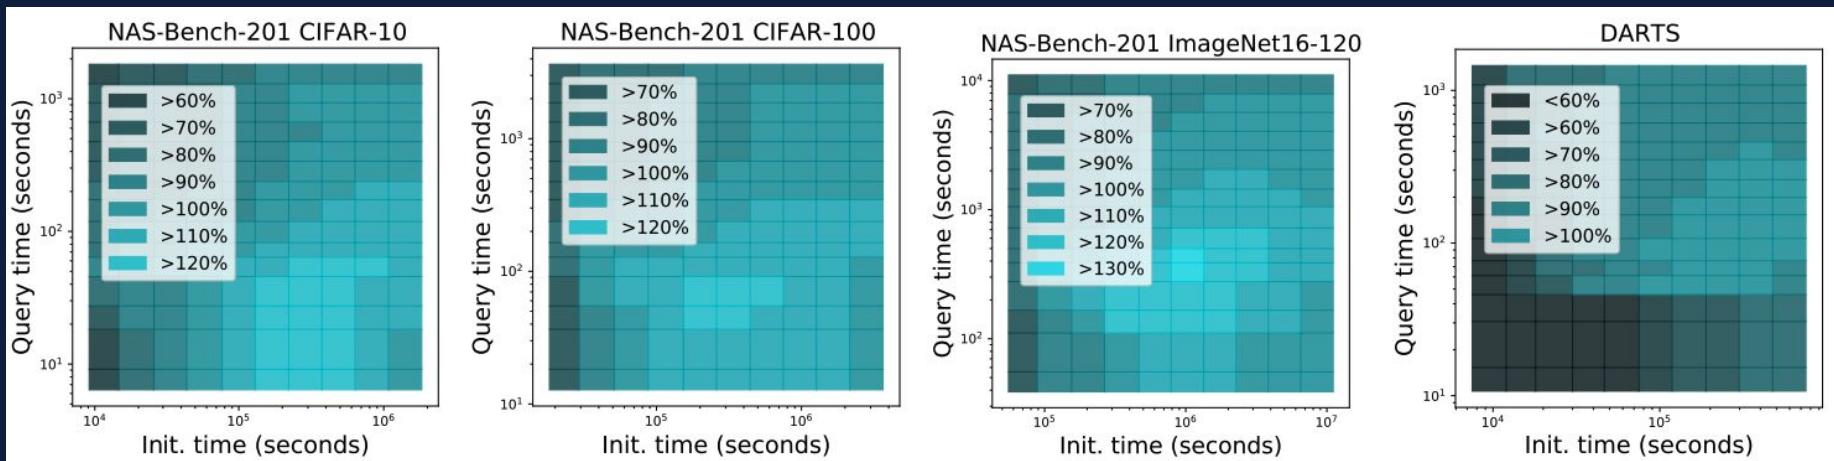
Performance predictors





OMNI: The Omnipotent Predictor

Combine best predictors from each family



Roadmap

- Motivation and Introduction
- Performance Prediction
 - BANANAS
 - Learning curve extrapolation
 - Zero-cost proxies
- **Recommender Systems**
- De-biasing models



Recommender Systems

neural architecture search

Advanced Machine Learning Day 3: Neural Architecture Search
27K views • 3 years ago
Microsoft Research

How do you search over architectures? View presentation slides and more at ...
CC

Dense Net | What Is a Markov Chain | Probabilistic Transitions | Hidden Markov Model | Sparse...

1:28:02

Notes on experiments

- Three axes of comparison: initialization time, query time, correlation / rank correlation metrics
- Official implementation whenever possible
- Train-test data drawn w.r.t.
- Light hyperparameter tuning
 - Leaves the playing field
 - Cross-validation is often used during NAS

14:56

How Powerful are Performance Predictors in Neural Architecture Search? (15 min video)
296 views • 1 year ago
Abacus AI

15 min video for the NeurIPS 2021 paper How Powerful are Performance Predictors in Neural Architecture Search? Colin White ...

14:56

Different Types of Performance Predictors | Overview of the Main Families of Performance...

11 min

How Powerful are Performance Predictors in Neural Architecture Search? (2 min video)
564 views • 1 year ago
Abacus AI

2 min video for the NeurIPS 2021 paper How Powerful are Performance Predictors in Neural Architecture Search? Colin White ...

2:01

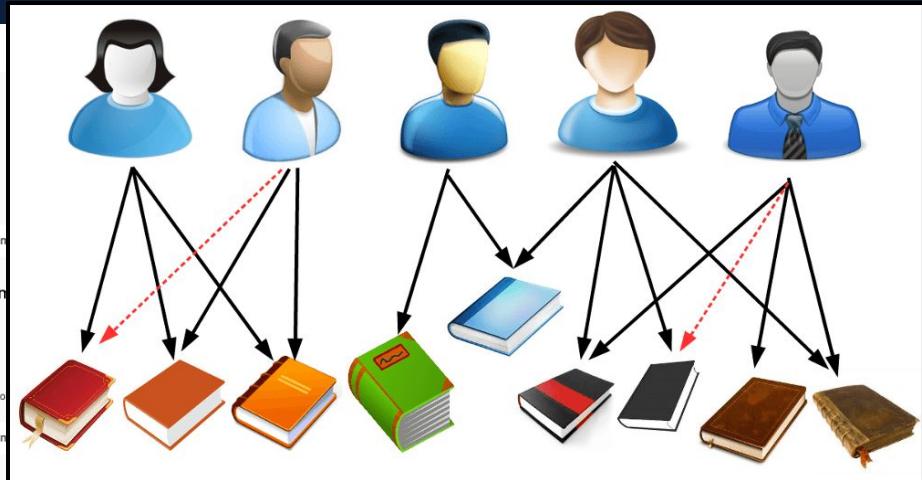
AutoML Fall School 21: Introduction to Neural Architecture Search
238 views • 3 months ago
AutoML Freiburg Hannover

Frank Hutter
University of Freiburg & Chair Center for AI
https://sites.google.com/view/automlsschool21/
@frankhutter, @AutoML_Hannover

BOSCH

These slides are available at <https://sites.google.com/view/automlsschool21/>

1:12:54



Frequently bought together

Total price: \$30.02

Add all three to Cart

Add all three to List

1.00

1.00

1.00

A Worrying Analysis of Recommender Systems

Are We Really Making Much Progress? A Worrying Analysis of Recent Neural Recommendation Approaches

Maurizio Ferrari Dacrema
Politecnico di Milano, Italy
maurizio.ferrari@polimi.it

Paolo Cremonesi
Politecnico di Milano, Italy
paolo.cremonesi@polimi.it

Dietmar Jannach
University of Klagenfurt, Austria
dietmar.jannach@aau.at

ABSTRACT

Deep learning techniques have become the method of choice for researchers working on algorithmic aspects of recommender systems. With the strongly increased interest in machine learning in general, it has, as a result, become difficult to keep track of what represents the state-of-the-art at the moment, e.g., for top-n recommendation tasks. At the same time, several recent publications point out problems in today's research practice in applied machine learning, e.g., in terms of the reproducibility of the results or the choice of the baselines when proposing new models.

In this work, we report the results of a systematic analysis of algorithmic proposals for top-n recommendation tasks. Specifically, we considered 18 algorithms that were presented at top-level research conferences in the last years. Only 7 of them could be reproduced with reasonable effort. For these methods, it however turned out that 6 of them can often be outperformed with comparatively simple heuristic methods, e.g., based on nearest-neighbor or graph-based techniques. The remaining one clearly outperformed the baselines but did not consistently outperform a well-tuned non-neural baseline. Overall, we conclude that

1 INTRODUCTION

Within only a few years, deep learning techniques have started to dominate the landscape of algorithmic research in recommender systems. Novel methods were proposed for a variety of settings and algorithmic tasks, including top-n recommendation based on long-term preference profiles or for session-based recommendation scenarios [36]. Given the increased interest in machine learning in general, the corresponding number of recent research publications, and the success of deep learning techniques in other fields like vision or language processing, one could expect that substantial progress resulted from these works also in the field of recommender systems. However, indications exist in other application areas of machine learning that the achieved progress—measured in terms of accuracy improvements over existing models—is not always as strong as expected.

Lin [25], for example, discusses two recent neural approaches in the field of information retrieval that were published at top-level conferences. His analysis reveals that the new methods do *not* significantly outperform existing baseline methods when these

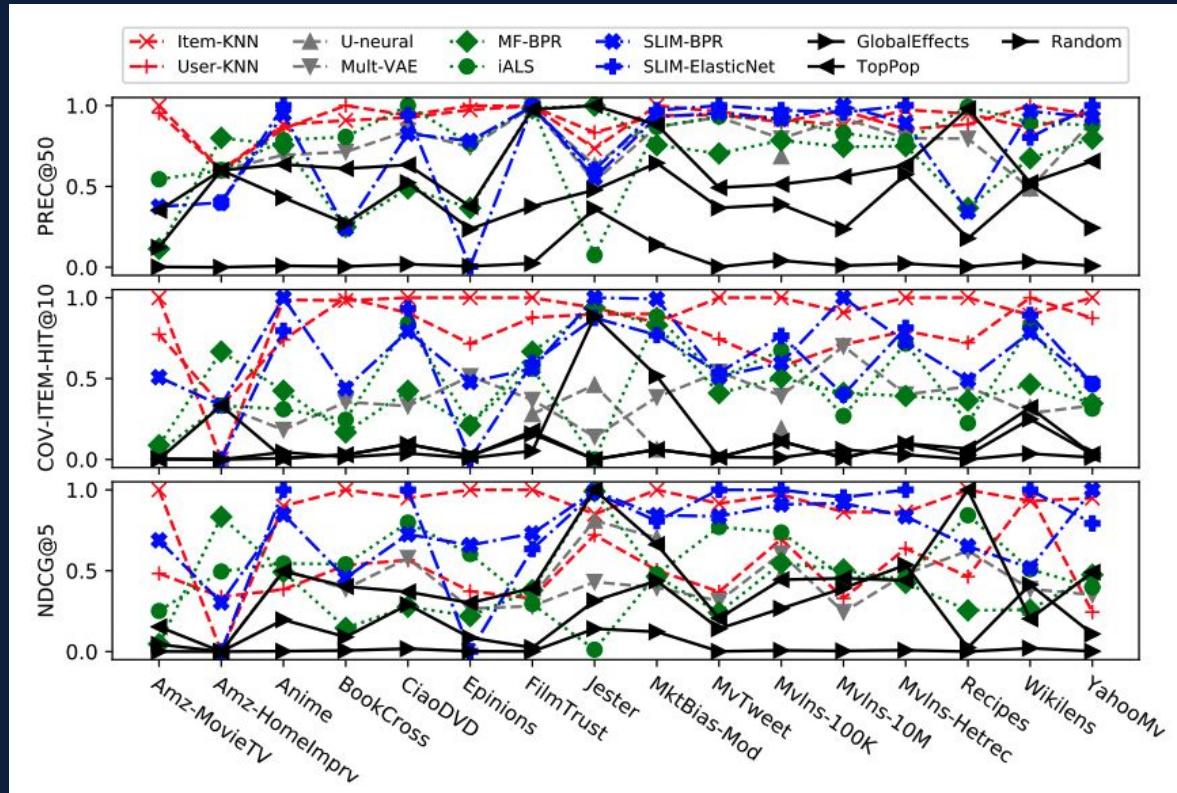
<i>Family</i>	<i>Method</i>	<i>Description</i>
Non-personalized	TopPopular	Recommends the most popular items to everyone [18]
Nearest-Neighbor	UserKNN	User-based k-nearest neighbors [58]
	ItemKNN	Item-based k-nearest neighbors [61]
Graph-based	$P^3\alpha$	A graph-based method based on random walks [16]
	$RP^3\beta$	An extension of $P^3\alpha$ [54]
Content-Based and Hybrid	ItemKNN-CBF	ItemKNN with content-based similarity [43]
	ItemKNN-CFCBF	A simple item-based hybrid CBF/CF approach [50]
	UserKNN-CBF	UserKNN with content-based similarity
	UserKNN-CFCBF	A simple user-based hybrid CBF/CF approach
Non-Neural Machine Learning	iALS	Matrix factorization for implicit feedback data [33]
	pureSVD	A basic matrix factorization method [18]
	SLIM	A scalable linear model [36, 52]
	EASE ^R	A recent linear model, similar to auto-encoders [63]

Meta-Learning for Recommender Systems

- 24 Algorithms, up to 100 hyperparameters, 85 datasets, 315 metrics

Rank	Item-KNN	P3alpha	SLIM-BPR	EASE-R	RP3beta	SVD	SLIM-ElasticNet	iALS	NMF	User-KNN	MF-Funk	TopPop	MF-Asy	MF-BPR	Mult-VAE	U-neural	GlobalEffects	CoClustering	Random	SlopeOne
Min.	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	2	1	9	7
Max.	14	18	14	18	17	16	17	19	14	17	18	19	16	17	20	20	20	19	20	20
Mean	2.3	4.2	4.7	5.3	6	6	7	7	7.1	7.6	9.4	10.4	10.7	11.2	11.7	12.3	13.3	14.9	16.2	16.7

Meta-Learning for Recommender Systems

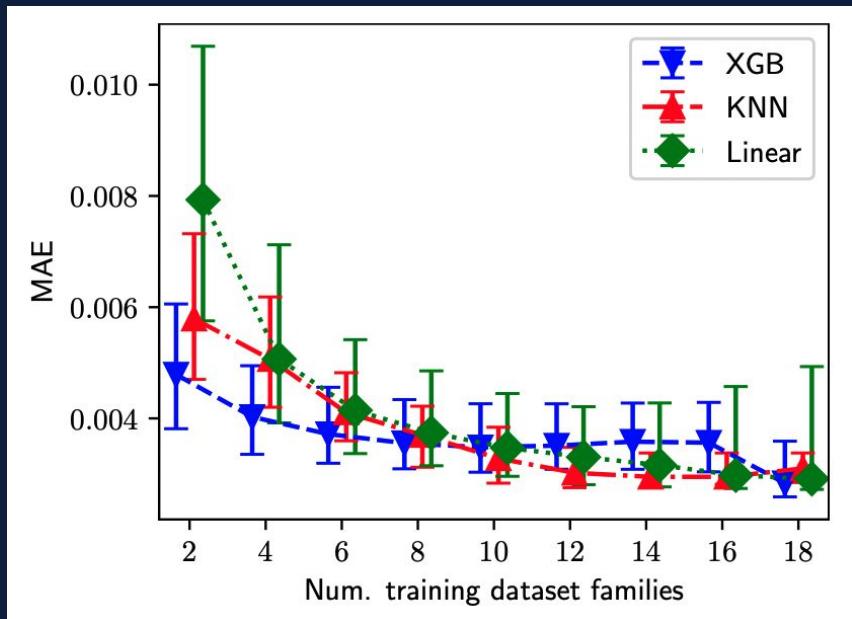




RECZILLA

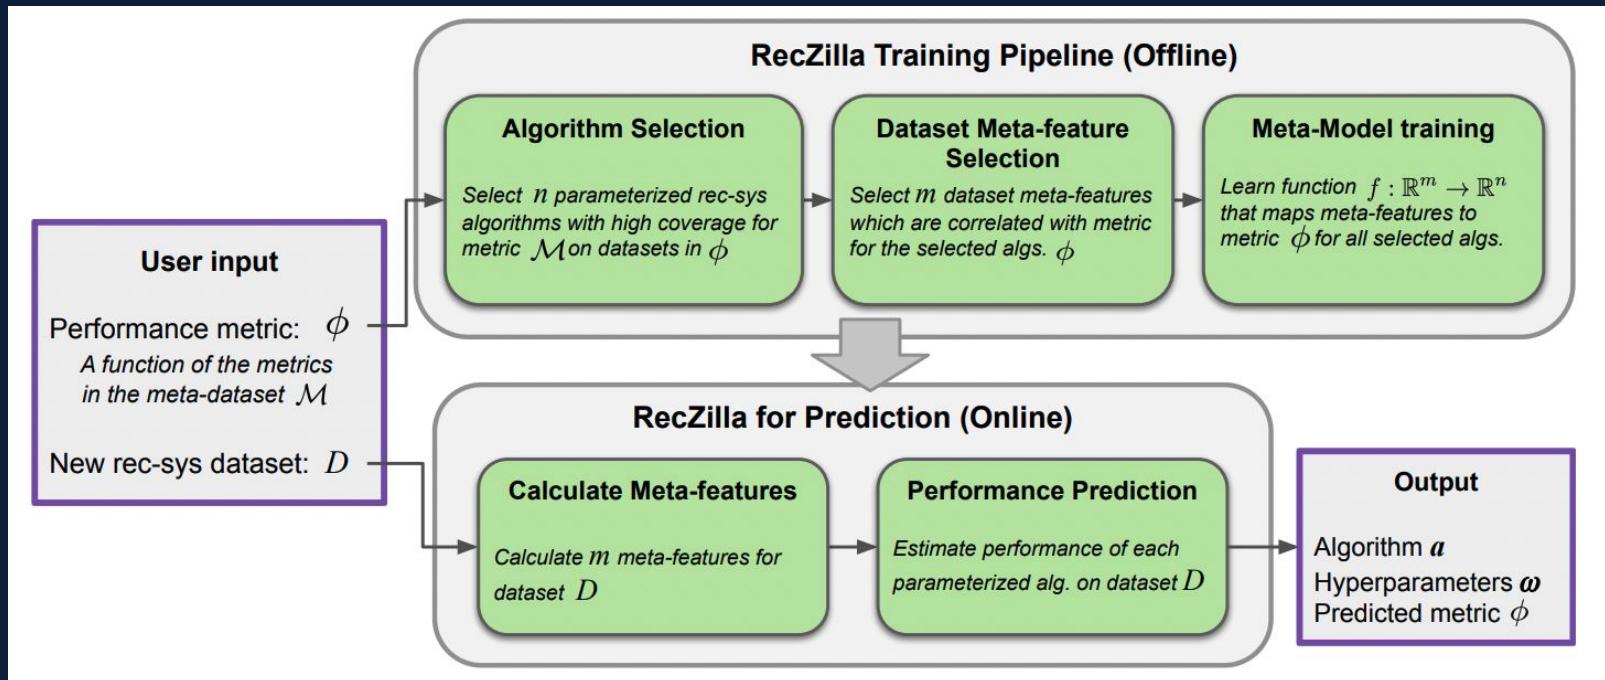
Dataset meta-features

- User distribution
- Item distribution
- Interaction distribution
- Landmarkers





RECZILLA



Roadmap

- Motivation and Introduction
- Performance Prediction
 - BANANAS
 - Learning curve extrapolation
 - Zero-cost proxies
- Recommender Systems
- De-biasing models



Face recognition

- 1. For one-to-one matching, the team saw higher rates of false positives for Asian and African American faces relative to images of Caucasians.** The differentials often ranged from a factor of 10 to 100 times, depending on the individual algorithm. False positives might present a security concern to the system owner, as they may allow access to impostors.
- 2. Among U.S.-developed algorithms, there were similar high rates of false positives in one-to-one matching for Asians, African Americans and native groups** (which include Native American, American Indian, Alaskan Indian and Pacific Islanders). The American Indian demographic had the highest rates of false positives.

'The Computer Got It Wrong': How Facial Recognition Led To False Arrest Of Black Man

June 24, 2020 · 8:00 AM ET



<https://www.nist.gov/news-events/news/2019/12/nist-study-evaluates-effects-race-age-sex-face-recognition-software>

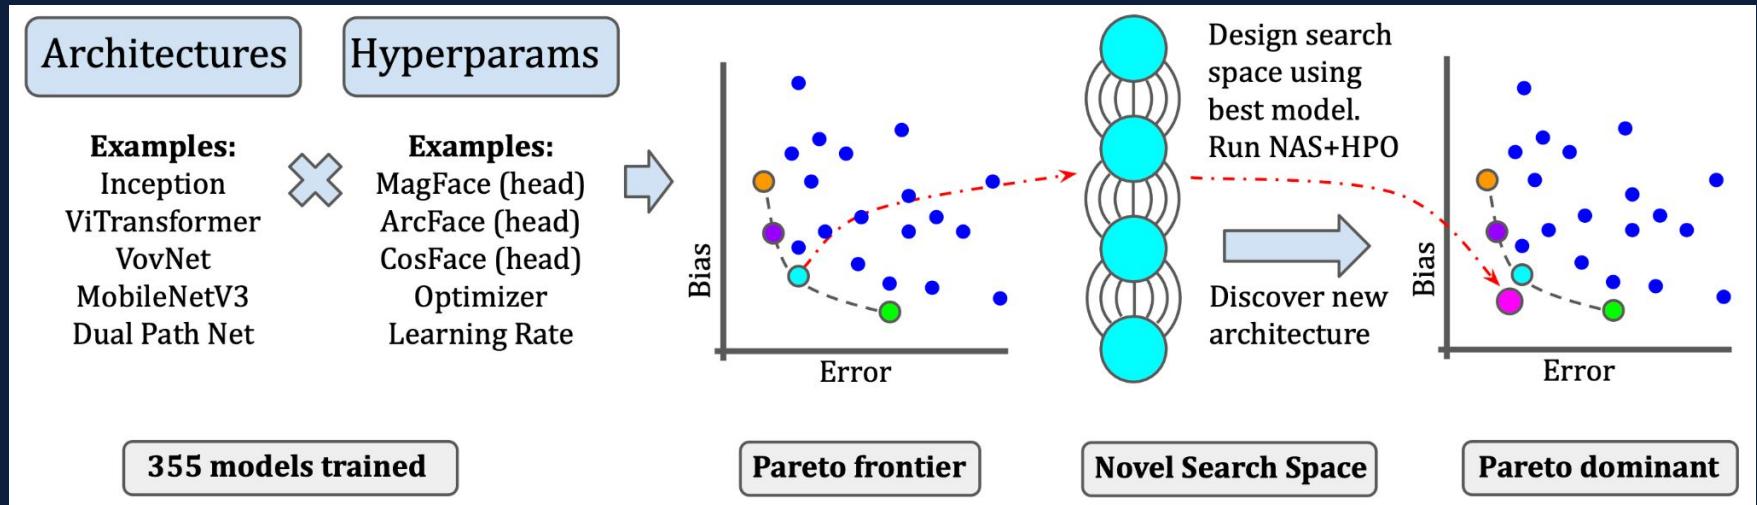
A US government study confirms most face recognition systems are racist

by Karen Hao December 20, 2019

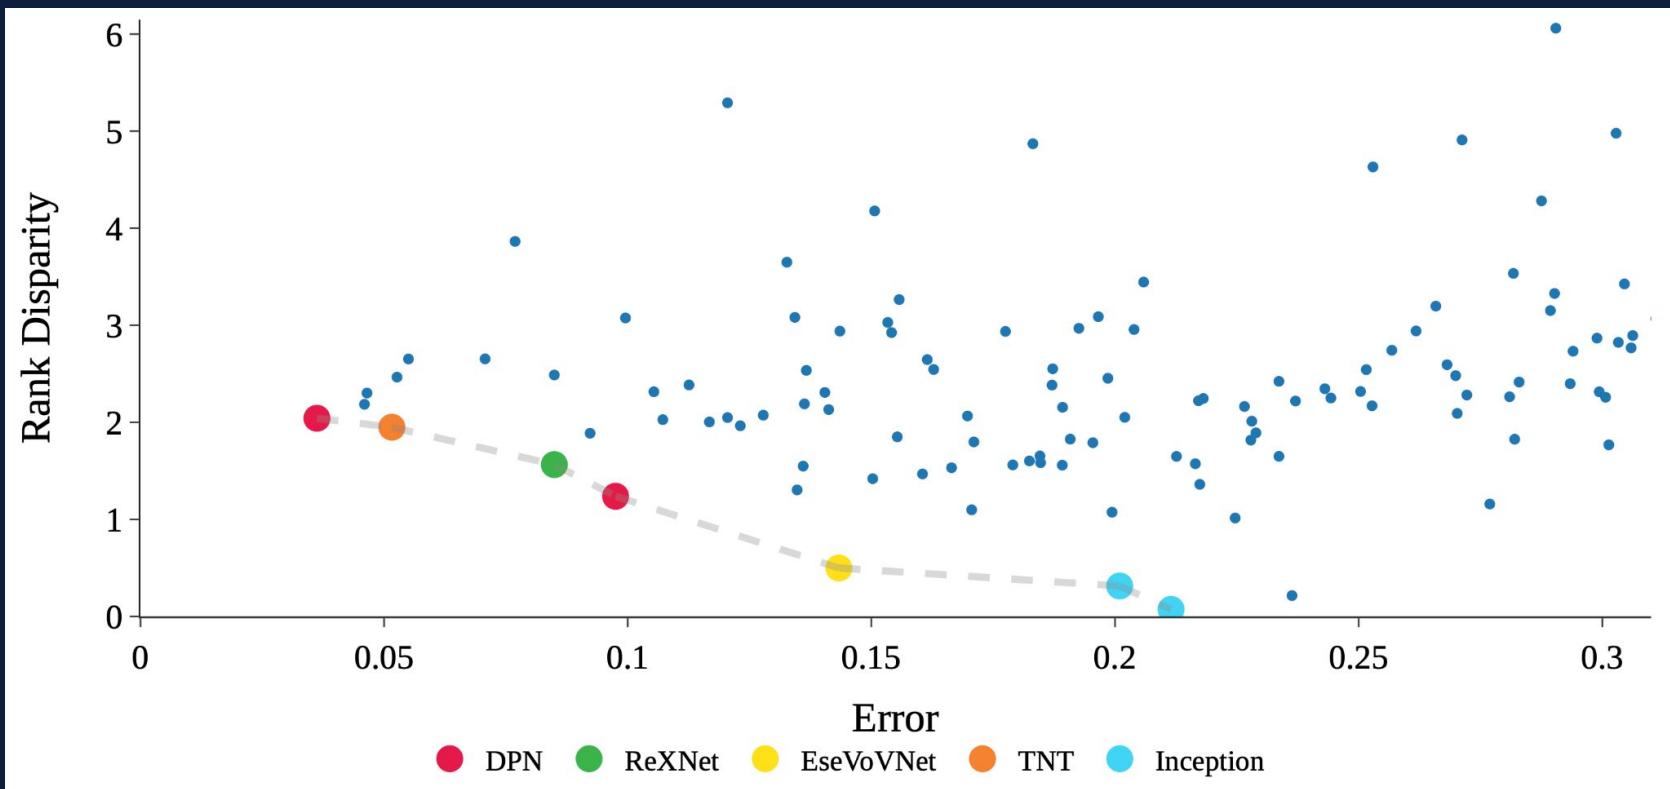


A U.S. Customs and Border Protection officer helps a passenger navigate a facial recognition kiosk at the airport.
DAVID J. PHILLIP/AP

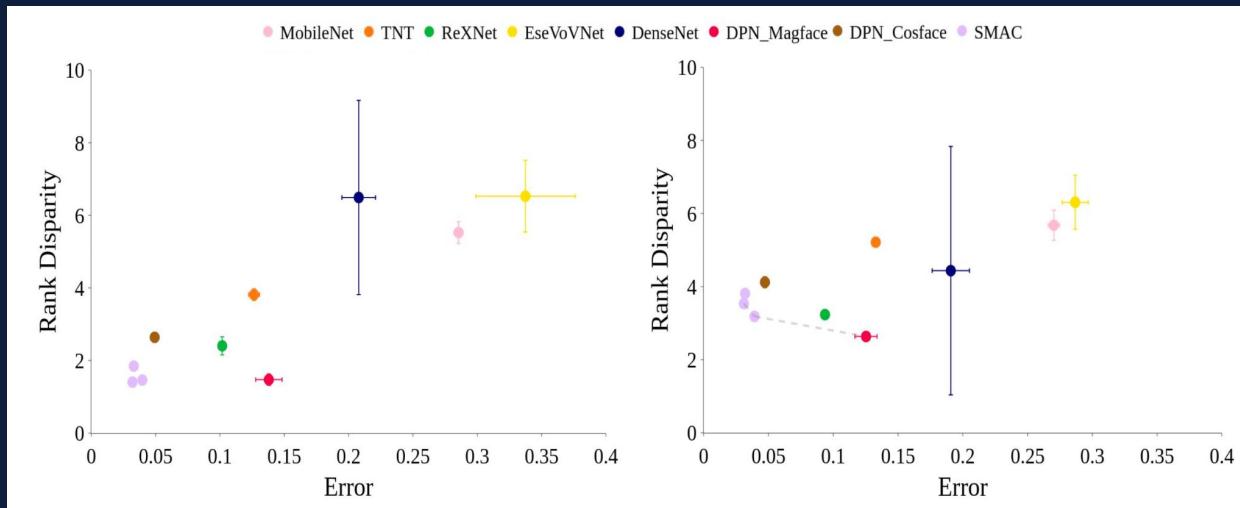
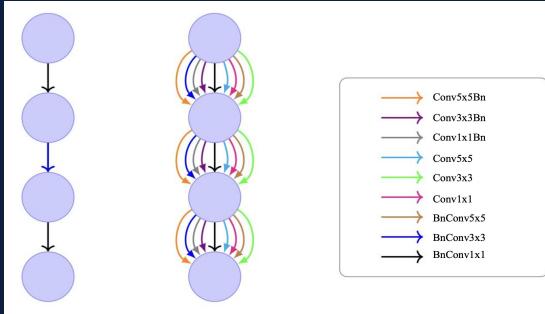
Face recognition



Face recognition



Face recognition



Thanks! Questions?



Colin White

colin@abacus.ai

<https://abacus.ai/publications>