


# An Introduction to Neural Architecture Search


Colin White, RealityEngines.AI

# Deep learning

- Explosion of interest since 2012
- Very powerful machine learning technique
- Huge variety of neural networks for different tasks
- Key ingredients took years to develop
- Algorithms are getting increasingly more specialized and complicated



# Deep learning




Algorithms are getting increasingly more specialized and complicated

E.g. accuracy on ImageNet has steadily improved over 10 years

# Neural architecture search

- What if an algorithm could do this for us?
- Neural architecture search (NAS) is a hot area of research
- Given a dataset, define a search space of architectures, then use a search strategy to find the best architecture for your dataset

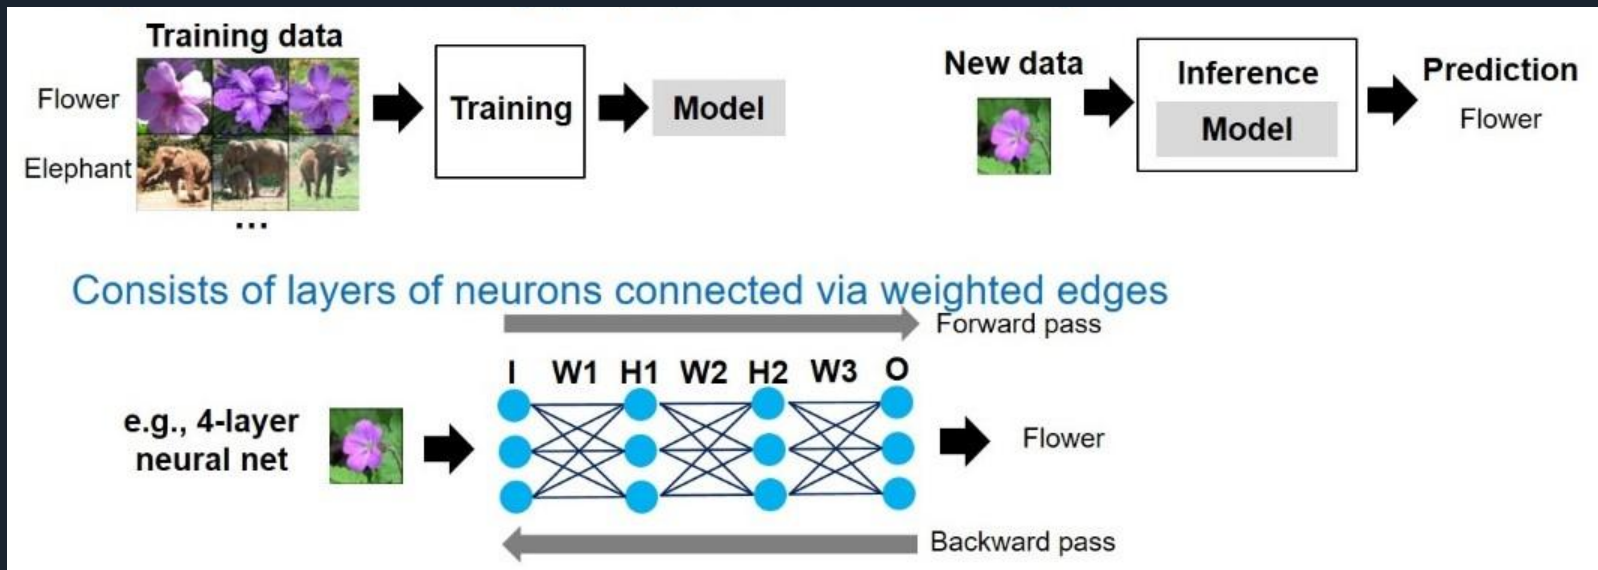


# Outline

- Introduction to NAS
- Background on deep learning
- Automated machine learning
- Optimization techniques
- NAS Framework
  - Search space
  - Search strategy
  - Evaluation method
- Conclusions

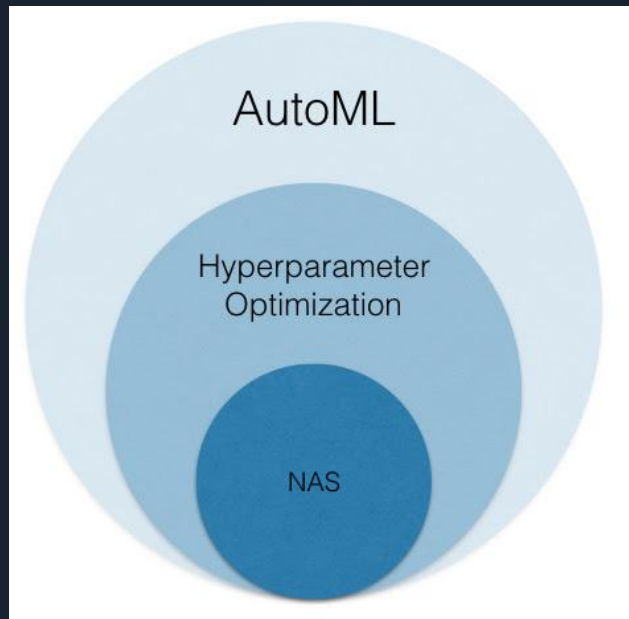
# Background - deep learning

- Studied since the 1940s - simulate the human brain
- “Neural networks are the second-best way to do almost anything” - JS Denker, 2000s
- Breakthrough: 2012 ImageNet competition [Krizhevsky, Sutskever, and Hinton]



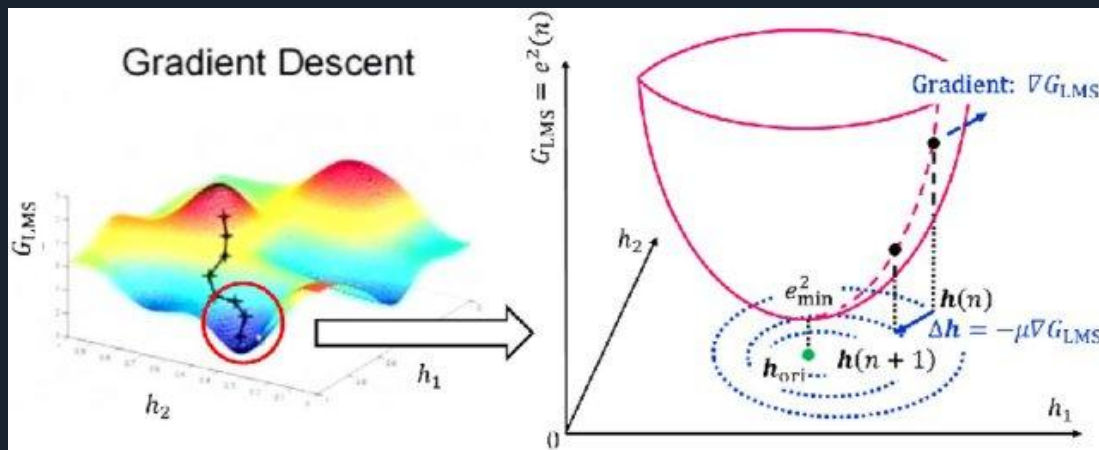
# Automated Machine Learning

- Automated machine learning
  - Data cleaning, model selection, HPO, NAS, ...
- Hyperparameter optimization (HPO)
  - Learning rate, dropout rate, batch size, ...
- Neural architecture search
  - Finding the best neural architecture



# Optimization


- Zero'th order optimization (used for HPO, NAS)
  - Bayesian Optimization
- First order optimization (used for Neural nets)
  - Gradient descent / stochastic gradient descent
- Second order optimization
  - Newton's method






# Zero'th order optimization


- Grid search
- Random search
- Bayesian Optimization
  - Use the results of the previous guesses to make the next guess



Grid Search



Random Search



Adaptive Selection

# Outline

- Introduction to NAS
- Background on deep learning
- Automated machine learning
- Optimization techniques
- NAS Framework
  - Search space
  - Search strategy
  - Evaluation method
- Conclusions

# Neural Architecture Search

## Search space

- Cell-based search space
- Macro vs micro search

## Search strategy

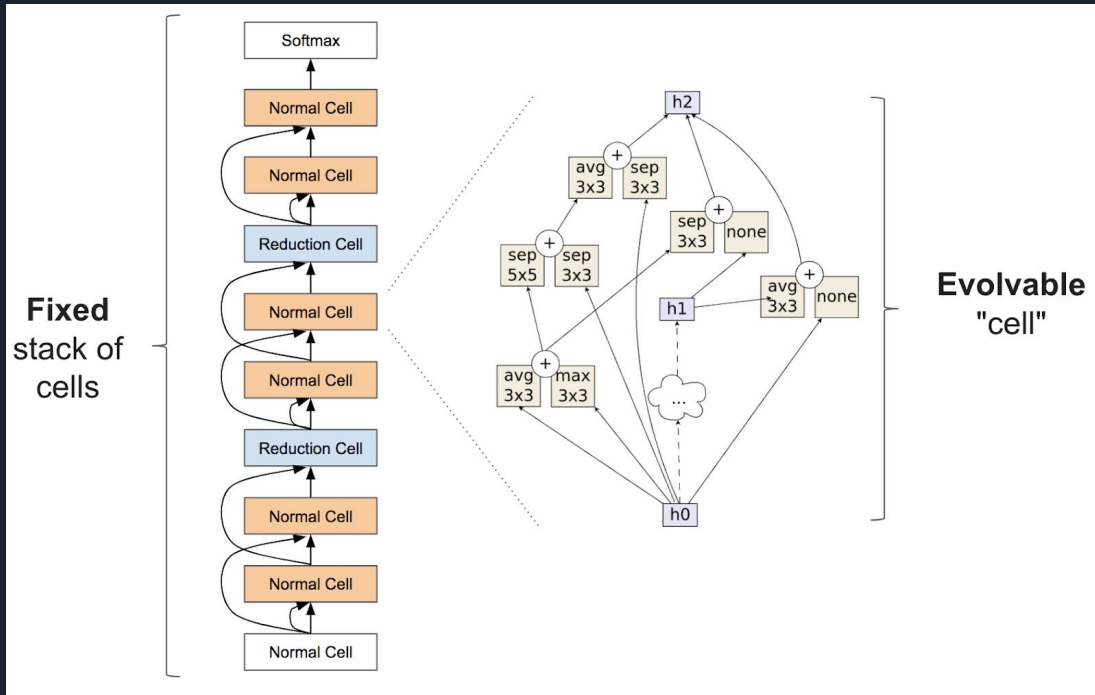
- Reinforcement learning
- Continuous methods
- Bayesian optimization
- Evolutionary algorithm

## Evaluation method

- Full training
- Partial training
- Training with shared weights

# Search space

- Macro vs micro search
- Progressive search
- Cell-based search space



[Zoph, Le '16]

# Bayesian optimization

- Popular method [Golovin et al. '17], [Jin et al. '18], [Kandasamy et al. '18]
- Great method to optimize an expensive function

- Fix a dataset (e.g. CIFAR-10, MNIST)
- Define a *search space*  $A$  (e.g., 20 layers of {conv, pool, ReLU, dense})
- Define an *objective function*  $f:A\rightarrow[0,1]$ 
  - $f(a)$  = validation accuracy of  $a$  after training
- Define a *distance function*  $d(a_1, a_2)$  between architectures
  - Quick to evaluate. If  $d(a_1, a_2)$  is small,  $|f(a_1) - f(a_2)|$  is small

# Bayesian optimization

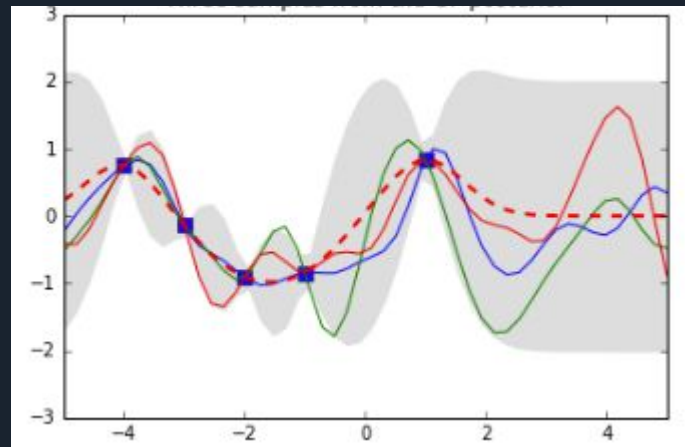
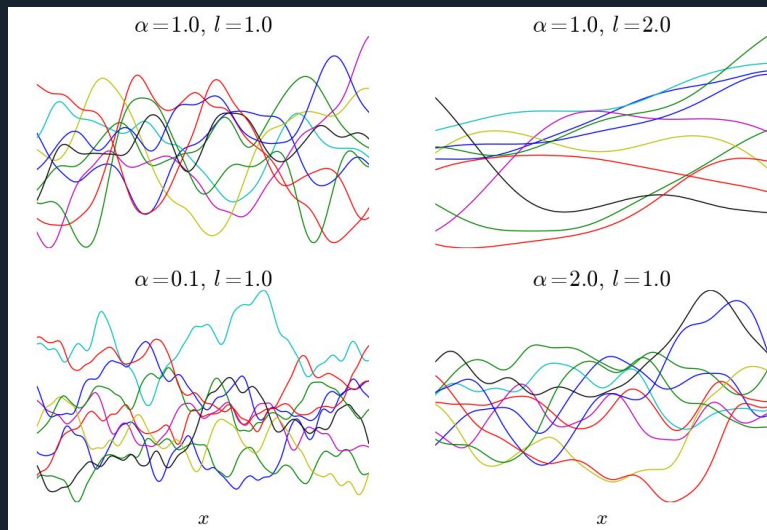
- Fix a dataset (e.g. CIFAR-10, MNIST)
- Define a *search space*  $A$  (e.g., 20 layers of {conv, pool, ReLU, dense})
- Define an *objective function*  $f:A\rightarrow[0,1]$ 
  - $f(a)$  = validation accuracy of  $a$  after training
- Define a *distance function*  $d(a_1, a_2)$  between architectures
  - Quick to evaluate. If  $d(a_1, a_2)$  is small,  $|f(a_1) - f(a_2)|$  is small

Goal: find  $a \in A$  which maximizes  $f(a)$

- Choose several random architectures  $a$  and evaluate  $f(a)$
- In each iteration  $i$ :
  - Use  $f(a_1) \dots f(a_{i-1})$  to choose new  $a_i$
  - Evaluate  $f(a_i)$

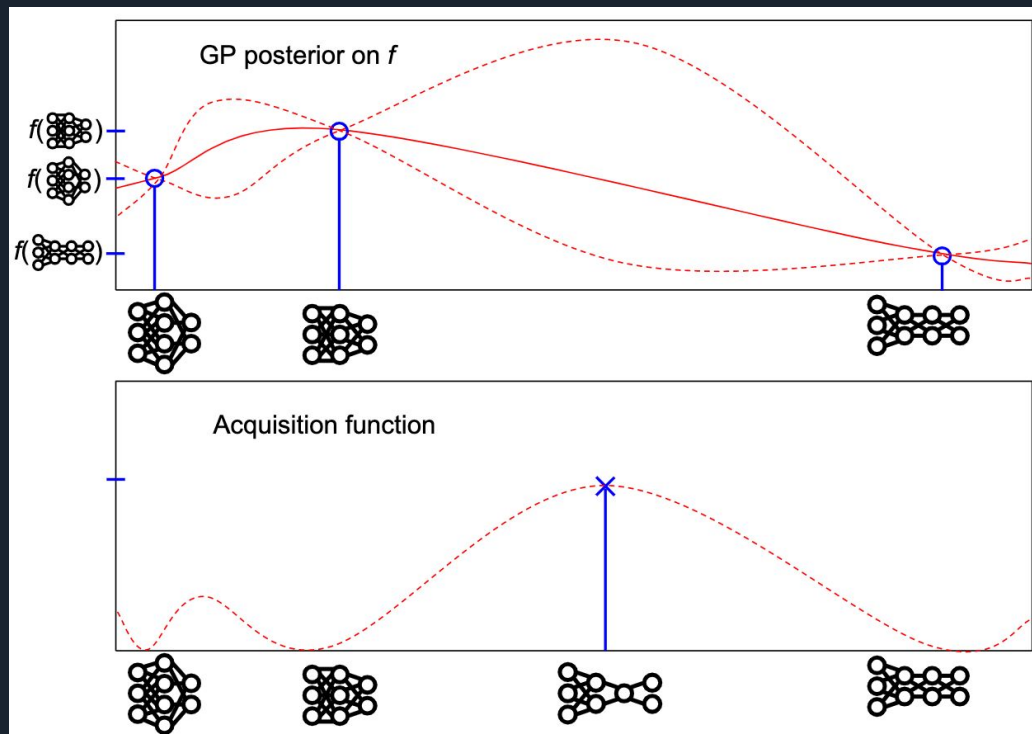
# Gaussian process

- Assume the distribution  $f(A)$  is *smooth*
- The deviations look like Gaussian noise
- Update as we get more information



# Acquisition function

- In each iteration, find the architecture with the largest ***expected improvement***







# DARTS: Differentiable Architecture Search

- Relax NAS to a *continuous* problem
- Use gradient descent (just like normal parameters)


[Liu et al. '18]




(a)



(b)



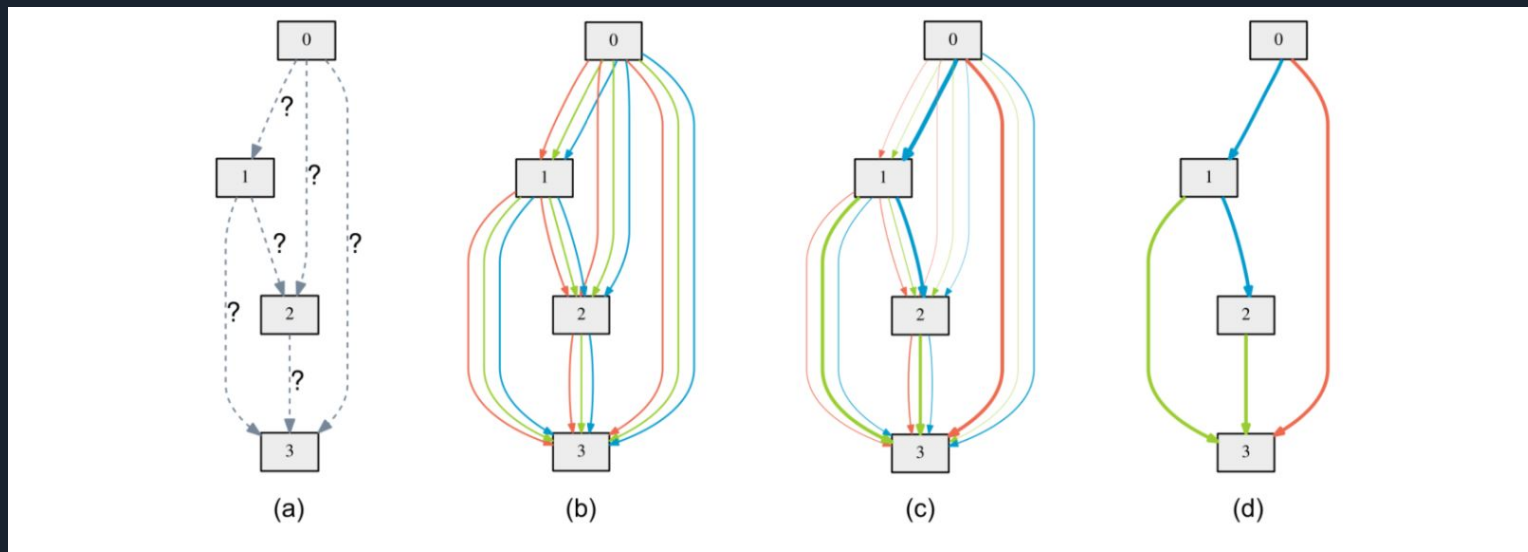
(c)



(d)

# DARTS: Differentiable Architecture Search

[Liu et al. '18]

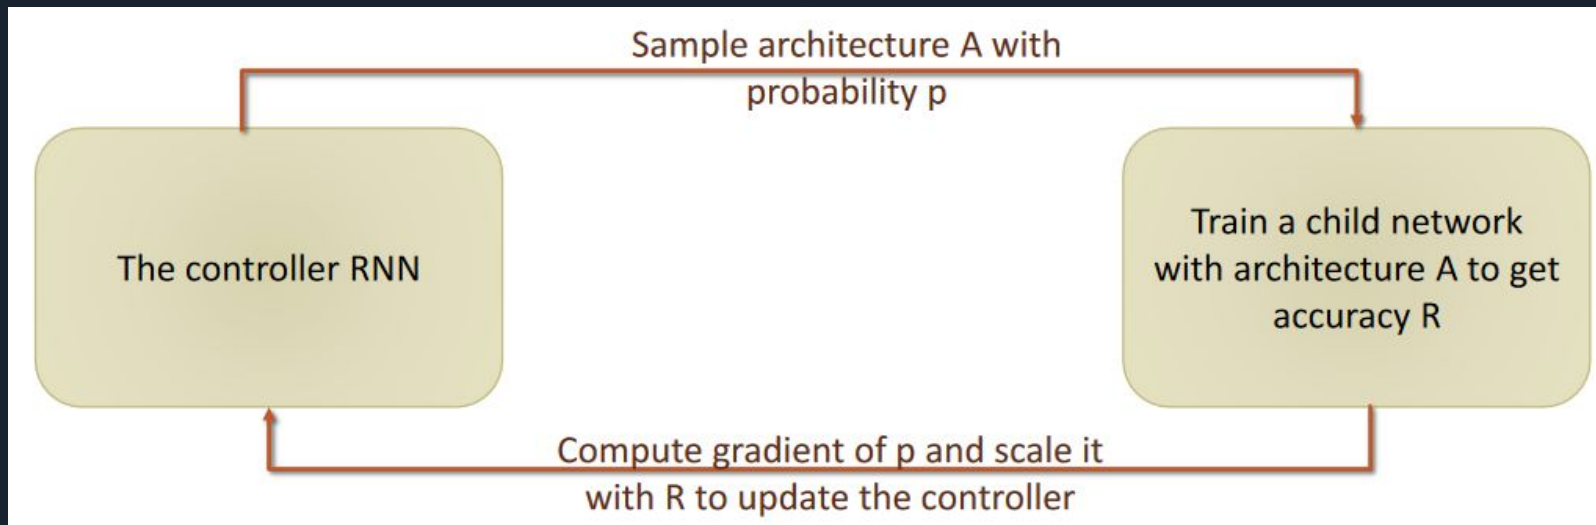


- Upsides: “one-shot”
- Downside: may only work in “micro” search setting

# Reinforcement Learning

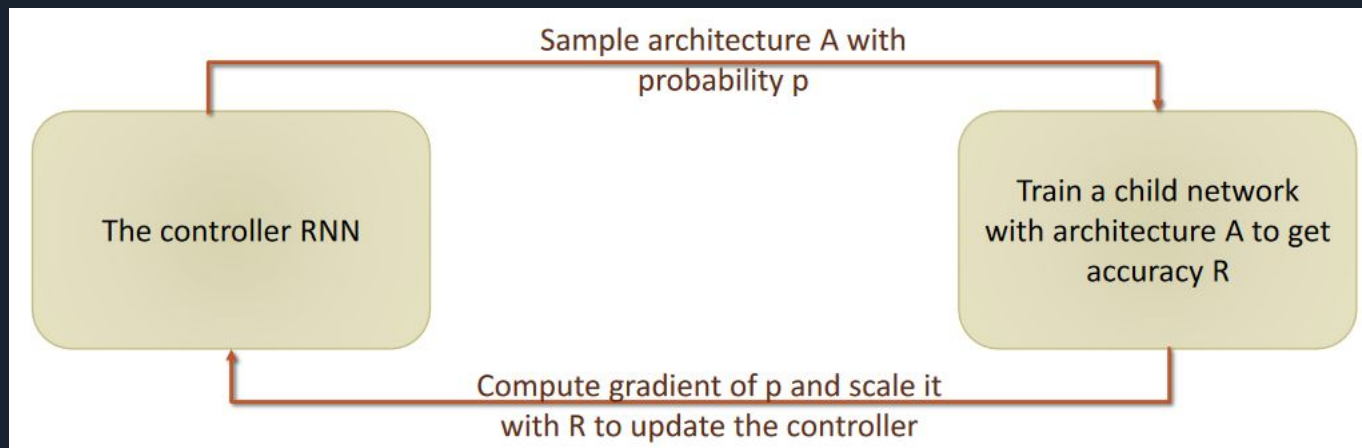
[Zoph, Le '16]

- **Controller** recurrent neural network
  - Chooses a new architecture in each round
  - Architecture is trained and evaluated
  - Controller receives feedback



# Reinforcement Learning

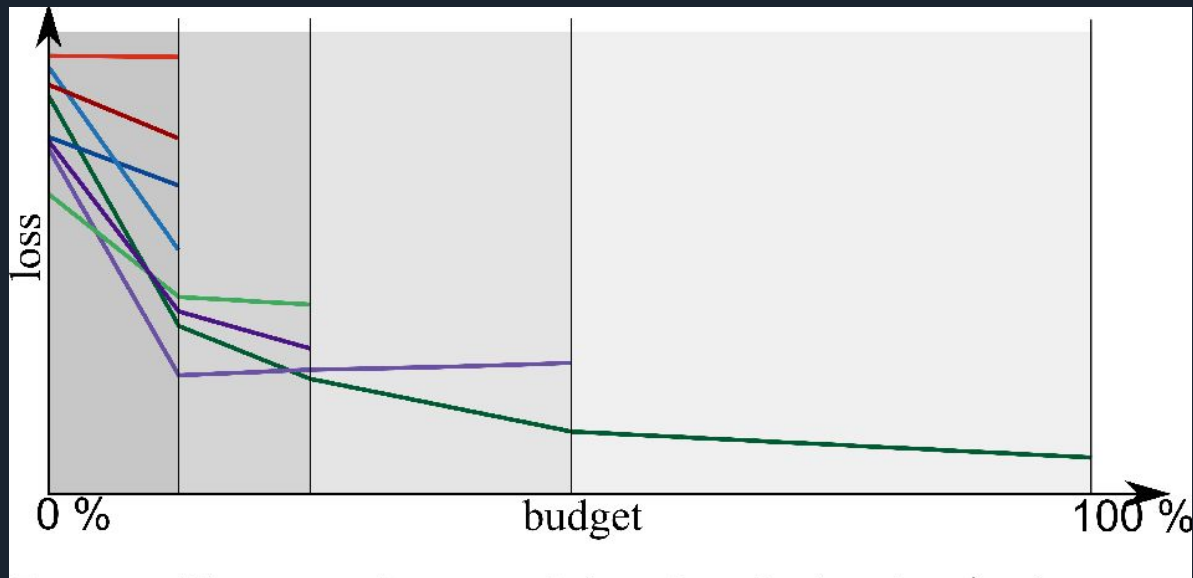
[Zoph, Le '16]



- Upside: much more powerful than BayesOpt, gradient descent
- Downsides: train a whole new network using neural networks; RL could be overkill

# Evaluation Strategy

- Full training
  - Simple
  - Accurate
- Partial training
  - Less computation
  - Less accurate
- Shared weights
  - Least computation
  - Least accurate



Source:

<https://www.automl.org/blog-2nd-automl-challenge/>

# Is NAS ready for widespread adoption?

- Hard to reproduce results [Li, Talwalkar '19]
  - Hard to compare different papers
  - Search spaces have been getting smaller
  - Random search is a strong baseline [Li, Talwalkar '19], [Sciuto et al. '19]
  - Recent papers are giving fair comparisons
  - NAS cannot yet consistently beat human engineers
- 
- Auto-Keras tutorial: <https://www.pyimagesearch.com/2019/01/07/auto-keras-and-automl-a-getting-started-guide/>
  - DARTS repo: <https://github.com/quark0/darts>

# Conclusion

- NAS: find the best neural architecture for a given dataset
- Search space, search strategy, evaluation method
- Search strategies: RL, BayesOpt, continuous optimization
- Not yet at the point of widespread adoption in industry

Thanks! Questions?